

STA380 Homework 2 Barton, Jace

Jace Barton

August 19, 2015

First, I load the libraries I will need.

```
library(RCurl)
## Warning: package 'RCurl' was built under R version 3.0.3
## Loading required package: bitops
library(ggplot2)
## Warning: package 'ggplot2' was built under R version 3.0.3
library(reshape)
## Warning: package 'reshape' was built under R version 3.0.3
library(plyr)
## Warning: package 'plyr' was built under R version 3.0.3
##
## Attaching package: 'plyr'
##
## The following objects are masked from 'package:reshape':
##
##      rename, round_any
library(tm)
## Warning: package 'tm' was built under R version 3.0.3
library(caret)
## Warning: package 'caret' was built under R version 3.0.3
## Loading required package: lattice
library(kknn)
## Warning: package 'kknn' was built under R version 3.0.3
##
## Attaching package: 'kknn'
##
## The following object is masked from 'package:caret':
##
##      contr.dummy
```

```

library(e1071)
## Warning: package 'e1071' was built under R version 3.0.3
library(randomForest)
## Warning: package 'randomForest' was built under R version 3.0.3
## randomForest 4.6-10
## Type rfNews() to see new features/changes/bug fixes.
library(arules)
## Warning: package 'arules' was built under R version 3.0.3
## Loading required package: Matrix
## Warning: package 'Matrix' was built under R version 3.0.3
##
## Attaching package: 'Matrix'
##
## The following object is masked from 'package:reshape':
##
##     expand
##
## The following objects are masked from 'package:base':
##
##     crossprod, tcrossprod
##
## Attaching package: 'arules'
##
## The following objects are masked from 'package:tm':
##
##     dissimilarity, inspect
##
## The following objects are masked from 'package:base':
##
##     %in%, write

```

I will run a random forest model in this homework. Thus, I will set the random seed so my results can be reproduced.

```
set.seed(722)
```

Now I'm ready to begin.

Graph Creation

Flights at ABIA

First, I will load in the data for the analysis.

```
AirportURLString =  
getURL("https://raw.githubusercontent.com/jacebarton/STA380/master/data/ABIA.  
csv", ssl.verifypeer=0L, followlocation = 1L)  
Airport = read.csv(text=AirportURLString)  
summary(Airport)
```

```
##      Year      Month      DayofMonth      DayOfWeek  
## Min.   :2008   Min.   : 1.00   Min.   : 1.00   Min.   :1.000  
## 1st Qu.:2008   1st Qu.: 3.00   1st Qu.: 8.00   1st Qu.:2.000  
## Median :2008   Median : 6.00   Median :16.00   Median :4.000  
## Mean   :2008   Mean   : 6.29   Mean   :15.73   Mean   :3.902  
## 3rd Qu.:2008   3rd Qu.: 9.00   3rd Qu.:23.00   3rd Qu.:6.000  
## Max.   :2008   Max.   :12.00   Max.   :31.00   Max.   :7.000  
##  
##      DepTime      CRSDepTime      ArrTime      CRSArrTime  
## Min.   :    1   Min.   : 55   Min.   :    1   Min.   :    5  
## 1st Qu.:  917   1st Qu.: 915   1st Qu.:1107   1st Qu.:1115  
## Median :1329   Median :1320   Median :1531   Median :1535  
## Mean   :1329   Mean   :1320   Mean   :1487   Mean   :1505  
## 3rd Qu.:1728   3rd Qu.:1720   3rd Qu.:1903   3rd Qu.:1902  
## Max.   :2400   Max.   :2346   Max.   :2400   Max.   :2400  
## NA's    :1413                NA's    :1567  
## UniqueCarrier      FlightNum      TailNum      ActualElapsedTime  
## WN      :34876   Min.   :    1                : 1104   Min.   : 22.0  
## AA      :19995   1st Qu.:  640   N678CA :  195   1st Qu.: 57.0  
## CO      : 9230   Median :1465   N511SW :  180   Median :125.0  
## YV      : 4994   Mean   :1917   N526SW :  176   Mean   :120.2  
## B6      : 4798   3rd Qu.:2653   N528SW :  172   3rd Qu.:164.0  
## XE      : 4618   Max.   :9741   N520SW :  168   Max.   :506.0  
## (Other):20749                (Other):97265   NA's    :1601  
## CRSElapsedTime      AirTime      ArrDelay      DepDelay  
## Min.   : 17.0   Min.   :  3.00   Min.   : -129.000   Min.   : -42.000  
## 1st Qu.: 58.0   1st Qu.: 38.00   1st Qu.:  -9.000   1st Qu.:  -4.000  
## Median :130.0   Median :105.00   Median :  -2.000   Median :   0.000  
## Mean   :122.1   Mean   : 99.81   Mean   :   7.065   Mean   :   9.171  
## 3rd Qu.:165.0   3rd Qu.:142.00   3rd Qu.:  10.000   3rd Qu.:   8.000  
## Max.   :320.0   Max.   :402.00   Max.   : 948.000   Max.   :875.000  
## NA's    :11     NA's    :1601   NA's    :1601     NA's    :1413  
##      Origin      Dest      Distance      TaxiIn  
## AUS      :49623   AUS      :49637   Min.   :   66   Min.   :  0.000  
## DAL      : 5583   DAL      : 5573   1st Qu.:  190   1st Qu.:  4.000  
## DFW      : 5508   DFW      : 5506   Median :  775   Median :  5.000  
## IAH      : 3704   IAH      : 3691   Mean   :  705   Mean   :  6.413  
## PHX      : 2786   PHX      : 2783   3rd Qu.:1085   3rd Qu.:  7.000
```

```

## DEN      : 2719   DEN      : 2673   Max.      :1770   Max.      :143.000
## (Other):29337   (Other):29397               NA's      :1567
##   TaxiOut      Cancelled      CancellationCode      Diverted
## Min.      : 1.00   Min.      :0.00000      :97840      Min.      :0.000000
## 1st Qu.: 9.00   1st Qu.:0.00000      A: 719      1st Qu.:0.000000
## Median :12.00   Median :0.00000      B: 605      Median :0.000000
## Mean    :13.96   Mean    :0.01431      C: 96       Mean    :0.001824
## 3rd Qu.:16.00   3rd Qu.:0.00000      3rd Qu.:0.000000
## Max.    :305.00   Max.    :1.00000      Max.    :1.000000
## NA's     :1419
##   CarrierDelay   WeatherDelay      NASDelay      SecurityDelay
## Min.      : 0.00   Min.      : 0.00   Min.      : 0.00   Min.      : 0.00
## 1st Qu.: 0.00   1st Qu.: 0.00   1st Qu.: 0.00   1st Qu.: 0.00
## Median : 0.00   Median : 0.00   Median : 2.00   Median : 0.00
## Mean    :15.39   Mean    : 2.24   Mean    :12.47   Mean    : 0.07
## 3rd Qu.:16.00   3rd Qu.: 0.00   3rd Qu.:16.00   3rd Qu.: 0.00
## Max.    :875.00   Max.    :412.00   Max.    :367.00   Max.    :199.00
## NA's     :79513   NA's     :79513   NA's     :79513   NA's     :79513
## LateAircraftDelay
## Min.      : 0.00
## 1st Qu.: 0.00
## Median : 6.00
## Mean    :22.97
## 3rd Qu.:30.00
## Max.    :458.00
## NA's     :79513

```

I immediately key in on delays as being the most interesting information in this dataset. For a given aircraft, are delays consistent for flights leaving Austin versus arriving in Austin?

To answer this, I first must split the dataset in two - one half for all of the flights leaving Austin, the other for all the flights arriving in Austin.

```

DepartureDelay <- data.frame(carrier=Airport$UniqueCarrier,
delay=Airport$DepDelay, leaving=Airport$Origin)
head(DepartureDelay)

```

```

##   carrier delay leaving
## 1      9E   345      MEM
## 2      AA    -5      AUS
## 3      YV     0      AUS
## 4      9E    -4      AUS
## 5      AA     1      AUS
## 6      NW    -9      AUS

```

This first step creates a data frame with all rows from the original data sets and columns for the airline, total delay for that flight, and the city from which the flight departed. I now want to filter this dataset to capture only Austin as the city of departure. I will also omit any rows where the departing city is unknown, as this means the flight was cancelled.

```
DepartureDelay = DepartureDelay[DepartureDelay$leaving == "AUS", ]
DepartureDelay = na.omit(DepartureDelay)
summary(DepartureDelay)
```

```
##      carrier      delay      leaving
## WN      :17343   Min.   :-36.000   AUS      :48893
## AA      : 9709   1st Qu.: -5.000   ABQ      :    0
## CO      : 4554   Median : -1.000   ATL      :    0
## YV      : 2456   Mean    :  7.425   BHM      :    0
## B6      : 2367   3rd Qu.:  5.000   BNA      :    0
## XE      : 2296   Max.    :875.000   BOS      :    0
## (Other):10168                (Other):    0
```

I can tell from the summary that this split the data almost exactly in half. I now need to aggregate delay information by carrier.

```
CarrierDepartureDelays = ddply(DepartureDelay, ~carrier, summarise,
mean=mean(delay), sd=sd(delay))
CarrierDepartureDelays
```

```
##      carrier      mean      sd
## 1         9E  3.656501 33.87182
## 2         AA  5.877536 28.25968
## 3         B6 10.451204 44.46459
## 4         CO  7.563900 32.73256
## 5         DL 12.099432 41.78871
## 6         EV 14.000000 40.72907
## 7         F9  1.599624 23.23526
## 8         MQ  7.820884 33.50115
## 9         NW  8.081967 48.06672
## 10        OH  9.926863 32.36067
## 11        OO  7.521761 33.37378
## 12        UA  5.833153 33.78858
## 13        US -0.778542 12.83104
## 14        WN  8.648158 24.97914
## 15        XE  5.597125 31.33525
## 16        YV  6.010586 35.25976
```

This completes my pre-processing for departing flights. I now need to do the same thing for arriving flights before final clean up.

```
ArrivalDelay <- data.frame(carrier=Airport$UniqueCarrier,
delay=Airport$ArrDelay, arriving=Airport$Dest)
head(DepartureDelay)
```

```
##      carrier delay leaving
## 2         AA    -5      AUS
## 3         YV     0      AUS
## 4         9E    -4      AUS
## 5         AA     1      AUS
## 6         NW    -9      AUS
## 7         CO    -9      AUS
```

```
ArrivalDelay = ArrivalDelay[ArrivalDelay$arriving == "AUS", ]
ArrivalDelay = na.omit(ArrivalDelay)
summary(ArrivalDelay)
```

```
##      carrier      delay      arriving
## WN      :17324   Min.   :-81.000   AUS      :48863
## AA      : 9708   1st Qu.: -9.000   ABQ      :    0
## CO      : 4555   Median : -1.000   ATL      :    0
## YV      : 2467   Mean    :  8.091   BNA      :    0
## B6      : 2365   3rd Qu.: 12.000   BOS      :    0
## XE      : 2288   Max.    :518.000   BWI      :    0
## (Other):10156                (Other):    0
```

```
CarrierArrivalDelays = ddply(ArrivalDelay, ~carrier, summarise,
mean=mean(delay), sd=sd(delay))
CarrierArrivalDelays
```

```
##      carrier      mean      sd
## 1         9E  3.518815 31.61852
## 2         AA  9.663473 34.46742
## 3         B6  9.610148 48.45161
## 4         CO  9.113063 35.08589
## 5         DL 12.979206 35.32323
## 6         EV 10.590571 38.63670
## 7         F9  5.172770 23.22650
## 8         MQ  6.428228 27.73061
## 9         NW 11.649123 49.13193
## 10        OH 15.274793 44.11506
## 11        OO  9.953854 36.11016
## 12        UA 12.237838 37.12183
## 13        US -2.640110 22.33975
## 14        WN  5.495324 29.89158
## 15        XE  6.173077 32.31526
## 16        YV 16.282529 47.60905
```

I now want to merge these two separate data frames. I also want to make the data more clear by using the airline name instead of the airline unique code.

```
CarrierDelays = merge(CarrierArrivalDelays, CarrierDepartureDelays,
by="carrier")
CarrierDelays$CarrierNames = c("Pinnacle", "American", "JetBlue",
"Continental", "Delta", "AtlanticSE", "Frontier", "Envoy", "Northwest",
"Comair", "SkyWest", "United", "US", "Southwest", "ExpressJet", "Mesa")
CarrierDelays
```

```
##      carrier      mean.x      sd.x      mean.y      sd.y CarrierNames
## 1         9E  3.518815 31.61852  3.656501 33.87182      Pinnacle
## 2         AA  9.663473 34.46742  5.877536 28.25968      American
## 3         B6  9.610148 48.45161 10.451204 44.46459      JetBlue
## 4         CO  9.113063 35.08589  7.563900 32.73256  Continental
## 5         DL 12.979206 35.32323 12.099432 41.78871      Delta
```

## 6	EV	10.590571	38.63670	14.000000	40.72907	AtlanticSE
## 7	F9	5.172770	23.22650	1.599624	23.23526	Frontier
## 8	MQ	6.428228	27.73061	7.820884	33.50115	Envoy
## 9	NW	11.649123	49.13193	8.081967	48.06672	Northwest
## 10	OH	15.274793	44.11506	9.926863	32.36067	Comair
## 11	OO	9.953854	36.11016	7.521761	33.37378	SkyWest
## 12	UA	12.237838	37.12183	5.833153	33.78858	United
## 13	US	-2.640110	22.33975	-0.778542	12.83104	US
## 14	WN	5.495324	29.89158	8.648158	24.97914	Southwest
## 15	XE	6.173077	32.31526	5.597125	31.33525	ExpressJet
## 16	YV	16.282529	47.60905	6.010586	35.25976	Mesa

Lastly, I will be interested in the total count of flights into and out of Austin for each airline. I will add this as a column. While I'm at it, I'll change the column names to be more meaningful.

```
CarrierDelays$Count = summary(Airport$UniqueCarrier)
colnames(CarrierDelays) = c("CarrierCode", "MeanDepartureDelay",
"SDDepartureDelay", "MeanArrivalDelay", "SDArrivalDelay", "CarrierNames",
"Count")
```

CarrierDelays


##	CarrierCode	MeanDepartureDelay	SDDepartureDelay	MeanArrivalDelay
## 1	9E	3.518815	31.61852	3.656501
## 2	AA	9.663473	34.46742	5.877536
## 3	B6	9.610148	48.45161	10.451204
## 4	CO	9.113063	35.08589	7.563900
## 5	DL	12.979206	35.32323	12.099432
## 6	EV	10.590571	38.63670	14.000000
## 7	F9	5.172770	23.22650	1.599624
## 8	MQ	6.428228	27.73061	7.820884
## 9	NW	11.649123	49.13193	8.081967
## 10	OH	15.274793	44.11506	9.926863
## 11	OO	9.953854	36.11016	7.521761
## 12	UA	12.237838	37.12183	5.833153
## 13	US	-2.640110	22.33975	-0.778542
## 14	WN	5.495324	29.89158	8.648158
## 15	XE	6.173077	32.31526	5.597125
## 16	YV	16.282529	47.60905	6.010586

##	SDArrivalDelay	CarrierNames	Count
## 1	33.87182	Pinnacle	2549
## 2	28.25968	American	19995
## 3	44.46459	JetBlue	4798
## 4	32.73256	Continental	9230
## 5	41.78871	Delta	2134
## 6	40.72907	AtlanticSE	825
## 7	23.23526	Frontier	2132
## 8	33.50115	Envoy	2663
## 9	48.06672	Northwest	121
## 10	32.36067	Comair	2986

```
## 11      33.37378      SkyWest  4015
## 12      33.78858      United   1866
## 13      12.83104        US     1458
## 14      24.97914    Southwest 34876
## 15      31.33525    ExpressJet 4618
## 16      35.25976      Mesa    4994
```

Now, what's all this been for? I want to get a picture of which airline is the best choice if I want to minimize my delays. I'd prefer my airline to be below average amongst all airlines in average delay on each leg of my trip - both departing and arriving. This can be seen in the following plot.

```
ggplot(CarrierDelays, aes(x=MeanArrivalDelay, y=MeanDepartureDelay,
label=CarrierNames)) + annotate("rect", xmin = -Inf, xmax =
mean(CarrierDelays$MeanArrivalDelay), ymin = -Inf, ymax =
mean(CarrierDelays$MeanDepartureDelay), fill= "green") +
  annotate("rect", xmin = -Inf, xmax = mean(CarrierDelays$MeanArrivalDelay),
ymin = mean(CarrierDelays$MeanDepartureDelay), ymax = Inf, fill= "yellow") +
  annotate("rect", xmin = mean(CarrierDelays$MeanArrivalDelay), xmax = Inf,
ymin = -Inf, ymax = mean(CarrierDelays$MeanDepartureDelay), fill= "yellow") +
  annotate("rect", xmin = mean(CarrierDelays$MeanArrivalDelay), xmax = Inf,
ymin = mean(CarrierDelays$MeanDepartureDelay), ymax = Inf, fill= "red") +
  geom_point(aes(size=CarrierDelays$Count)) +
  geom_vline(xintercept=mean(CarrierDelays$MeanArrivalDelay)) +
  geom_hline(yintercept=mean(CarrierDelays$MeanDepartureDelay)) +
  scale_size_continuous(range=c(3,15)) + geom_text(size=3, hjust=1)
```



Dots in the green square have smaller delays on average both arriving and departing, whereas dots in the red square have larger delays on average in both directions. The dot size is proportional to the number of flights the airline has into and out of Austin.

This graph can help a traveler determine which airline to take. For instance, if I want an airline with a lot of flights and good performance in getting me home on time, I'll choose American. On the other hand, if I want an airline with a lot of flights and I care most about getting to my non-Austin destination on time, I'll choose Southwest.

Text Analysis

Author Attribution

I am given approximately 50 New York Times articles each of 50 different authors as a training set, and another 50 articles each of the same authors as a test set. Can I accurately predict which author a given article from the test set belongs to?

For simplicity, this analysis will ignore words from the test data set which are not in the training data set.

First, I pass in a function I will need to read the text data.

```
readerPlain = function(fname){  
  readPlain(elem=list(content=readLines(fname)),  
               id=fname, language='en') }
```

Now, I will bring in the test data, using the procedure from the example in class.

```
#Get all files  
train_author_dirs = Sys.glob('../data/ReutersC50/C50train/*')  
train_file_list = NULL  
train_labels = NULL  
#build a single corpus  
for(author in train_author_dirs) {  
  author_name = substring(author, first=29)  
  files_to_add = Sys.glob(paste0(author, '/*.txt'))  
  train_file_list = append(train_file_list, files_to_add)  
  train_labels = append(train_labels, rep(author_name, length(files_to_add)))  
}  
  
# Need a more clever regex to get better names here  
train_all_docs = lapply(train_file_list, readerPlain)  
names(train_all_docs) = sub('.txt', '', names(train_all_docs))  
  
train_corpus = Corpus(VectorSource(train_all_docs))  
names(train_corpus) = train_file_list  
  
# Clean up tokens in corpus  
train_corpus = tm_map(train_corpus, tolower) # make everything lowercase
```

```

train_corpus = tm_map(train_corpus, removeNumbers) # remove numbers
train_corpus = tm_map(train_corpus, removePunctuation) # remove punctuation
train_corpus = tm_map(train_corpus, stripWhitespace) ## remove excess white-
space
train_corpus = tm_map(train_corpus, removeWords, stopwords("SMART"))

Train_Document_Term_Matrix = DocumentTermMatrix(train_corpus)
Train_Document_Term_Matrix # some basic summary statistics

## A document-term matrix (2500 documents, 31423 terms)
##
## Non-/sparse entries: 425955/78131545
## Sparsity           : 99%
## Maximal term length: 36
## Weighting          : term frequency (tf)

Train_Document_Term_Matrix = removeSparseTerms(Train_Document_Term_Matrix,
0.975)
tm::inspect(Train_Document_Term_Matrix[1:10,1:5])

## A document-term matrix (10 documents, 5 terms)
##
## Non-/sparse entries: 3/47
## Sparsity           : 94%
## Maximal term length: 10
## Weighting          : term frequency (tf)
##
##      Terms
## Docs ability abroad access account accounting
## 1      0      0      1      0      0
## 2      0      0      0      0      0
## 3      0      0      2      0      0
## 4      0      0      0      0      0
## 5      0      0      0      0      0
## 6      0      0      0      0      0
## 7      0      0      0      0      0
## 8      0      0      0      0      0
## 9      0      0      0      0      0
## 10     0      0      4      0      0

# Now a dense matrix
Train_Matrix = as.matrix(Train_Document_Term_Matrix)

```

I will repeat the above steps to build out the testing data.

```

test_author_dirs = Sys.glob('../data/ReutersC50/C50test/*')
test_file_list = NULL
test_labels = NULL
#build a single corpus
for(author in test_author_dirs) {
  author_name = substring(author, first=28)

```

```

files_to_add = Sys.glob(paste0(author, '/*.txt'))
test_file_list = append(test_file_list, files_to_add)
test_labels = append(test_labels, rep(author_name, length(files_to_add)))
}

# Need a more clever regex to get better names here
test_all_docs = lapply(test_file_list, readerPlain)
names(test_all_docs) = sub('.txt', '', names(test_all_docs))

test_corpus = Corpus(VectorSource(test_all_docs))
names(test_corpus) = test_file_list

# Clean up tokens in corpus
test_corpus = tm_map(test_corpus, tolower) # make everything lowercase
test_corpus = tm_map(test_corpus, removeNumbers) # remove numbers
test_corpus = tm_map(test_corpus, removePunctuation) # remove punctuation
test_corpus = tm_map(test_corpus, stripWhitespace) ## remove excess white-
space
test_corpus = tm_map(test_corpus, removeWords, stopwords("SMART"))

Test_Document_Term_Matrix = DocumentTermMatrix(test_corpus, control =
list(dictionary=Terms(Train_Document_Term_Matrix)) )
Test_Document_Term_Matrix # some basic summary statistics

## A document-term matrix (2500 documents, 1389 terms)
##
## Non-/sparse entries: 246565/3225935
## Sparsity          : 93%
## Maximal term length: 18
## Weighting          : term frequency (tf)

tm::inspect(Test_Document_Term_Matrix[1:10,1:5])

## A document-term matrix (10 documents, 5 terms)
##
## Non-/sparse entries: 9/41
## Sparsity          : 82%
## Maximal term length: 10
## Weighting          : term frequency (tf)
##
##      Terms
## Docs ability abroad access account accounting
## 1          0      0      0          1          3
## 2          0      0      0          0          0
## 3          0      0      3          0          0
## 4          1      1      0          0          0
## 5          0      0      0          0          0
## 6          0      0      0          0          0
## 7          0      0      0          0          0
## 8          1      0      4          0          0

```

```
##      9      1      0      4      0      0
##     10      0      0      0      0      0

# Now a dense matrix
Test_Matrix = as.matrix(Test_Document_Term_Matrix)
```

Now, I will build a Naive Bayes model to attempt to classify which test articles belong to which authors.

```
NaiveBayesModel = naiveBayes(Train_Matrix, as.factor(train_labels),
                             laplace=1)

NaiveBayesPredict = predict(object=NaiveBayesModel, newdata = Test_Matrix)
```

Now that the model is built, I can begin to look at results in different ways.

```
NaiveBayesResults = as.data.frame(table(NaiveBayesPredict, test_labels))
NaiveBayesResultsTable = table(NaiveBayesPredict, test_labels)
total_correct_vector = rep(0, 50)
predicted_per_author = rep(0, 50)
for (i in 1:length(NaiveBayesResultsTable[1,])){
  total_correct_vector[i] = NaiveBayesResultsTable[i, i]
  predicted_per_author[i] = sum(NaiveBayesResultsTable[i,])
}
```

```
CorrectByAuthor = data.frame(row.names(NaiveBayesResultsTable),
                             total_correct_vector)
CorrectByAuthor
```

```
##      row.names.NaiveBayesResultsTable. total_correct_vector
## 1                      AaronPressman                29
## 2                      AlanCrosby                   43
## 3                      AlexanderSmith                0
## 4                      BenjaminKangLim              16
## 5                      BernardHickey                8
## 6                      BradDorfman                 0
## 7                      DarrenSchuettler             0
## 8                      DavidLawder                 21
## 9                      EdnaFernandes               0
## 10                     EricAuchard                 1
## 11                     FumikoFujisaki              7
## 12                     GrahamEarnshaw              1
## 13                     HeatherScoffield             2
## 14                     JaneMacartney               0
## 15                     JanLopatka                  6
## 16                     JimGilchrist               39
## 17                     JoeOrtiz                   1
## 18                     JohnMastrini                2
## 19                     JonathanBirt                0
## 20                     JoWinterbottom             18
## 21                     KarlPenhaul                 2
```

## 22	KeithWeir	1
## 23	KevinDrawbaugh	0
## 24	KevinMorrison	1
## 25	KirstinRidley	0
## 26	KouroshKarimkhany	40
## 27	LydiaZajc	47
## 28	LynneO'Donnell	23
## 29	LynnleyBrowning	4
## 30	MarcelMichelson	12
## 31	MarkBendeich	0
## 32	MartinWolk	0
## 33	MatthewBunce	5
## 34	MichaelConnor	1
## 35	MureDickie	0
## 36	NickLouth	5
## 37	PatriciaCommins	0
## 38	PeterHumphrey	32
## 39	PierreTran	2
## 40	RobinSidel	20
## 41	RogerFillion	35
## 42	SamuelPerry	1
## 43	SarahDavison	0
## 44	ScottHillis	0
## 45	SimonCowell	0
## 46	TanEeLyn	0
## 47	TheresePoletti	6
## 48	TimFarrand	13
## 49	ToddNissen	0
## 50	WilliamKazer	0

```

PredictedForAuthor = data.frame(row.names(NaiveBayesResultsTable),
predicted_per_author)
SortedPredictedForAuthor = PredictedForAuthor[order(-
PredictedForAuthor$predicted_per_author),]
SortedPredictedForAuthor

```

##	row.names.NaiveBayesResultsTable.	predicted_per_author
## 2	AlanCrosby	642
## 27	LydiaZajc	593
## 8	DavidLawder	354
## 26	KouroshKarimkhany	243
## 16	JimGilchrist	139
## 41	RogerFillion	80
## 4	BenjaminKangLim	75
## 38	PeterHumphrey	74
## 1	AaronPressman	59
## 48	TimFarrand	50
## 28	LynneO'Donnell	29
## 15	JanLopatka	28
## 20	JoWinterbottom	24

## 40	RobinSidel	21
## 30	MarcelMichelson	20
## 47	TheresePoletti	12
## 5	BernardHickey	11
## 11	FumikoFujisaki	8
## 36	NickLouth	8
## 33	MatthewBunce	5
## 13	HeatherScofield	4
## 29	LynnleyBrowning	4
## 21	KarlPenhaul	3
## 18	JohnMastrini	2
## 22	KeithWeir	2
## 34	MichaelConnor	2
## 39	PierreTran	2
## 10	EricAuchard	1
## 12	GrahamEarnshaw	1
## 17	JoeOrtiz	1
## 24	KevinMorrison	1
## 42	SamuelPerry	1
## 49	ToddNissen	1
## 3	AlexanderSmith	0
## 6	BradDorfman	0
## 7	DarrenSchuettler	0
## 9	EdnaFernandes	0
## 14	JaneMacartney	0
## 19	JonathanBirt	0
## 23	KevinDrawbaugh	0
## 25	KirstinRidley	0
## 31	MarkBendeich	0
## 32	MartinWolk	0
## 35	MureDickie	0
## 37	PatriciaCommins	0
## 43	SarahDavison	0
## 44	ScottHillis	0
## 45	SimonCowell	0
## 46	TanEeLyn	0
## 50	WilliamKazer	0

```
OverallClassificationRate = sum(total_correct_vector)/2500
OverallClassificationRate
```

```
## [1] 0.1776
```

```
PrecisionRateByAuthor = data.frame(row.names(NaiveBayesResultsTable),
total_correct_vector/predicted_per_author)
PrecisionRateByAuthor
```

```
## row.names.NaiveBayesResultsTable.
## 1 AaronPressman
## 2 AlanCrosby
## 3 AlexanderSmith
```

```

## 4 BenjaminKangLim
## 5 BernardHickey
## 6 BradDorfman
## 7 DarrenSchuettler
## 8 DavidLawder
## 9 EdnaFernandes
## 10 EricAuchard
## 11 FumikoFujisaki
## 12 GrahamEarnshaw
## 13 HeatherScoffield
## 14 JaneMacartney
## 15 JanLopatka
## 16 JimGilchrist
## 17 JoeOrtiz
## 18 JohnMastrini
## 19 JonathanBirt
## 20 JoWinterbottom
## 21 KarlPenhaul
## 22 KeithWeir
## 23 KevinDrawbaugh
## 24 KevinMorrison
## 25 KirstinRidley
## 26 KouroshKarimkhany
## 27 LydiaZajc
## 28 LynneO'Donnell
## 29 LynnleyBrowning
## 30 MarcelMichelson
## 31 MarkBendeich
## 32 MartinWolk
## 33 MatthewBunce
## 34 MichaelConnor
## 35 MureDickie
## 36 NickLouth
## 37 PatriciaCommins
## 38 PeterHumphrey
## 39 PierreTran
## 40 RobinSidel
## 41 RogerFillion
## 42 SamuelPerry
## 43 SarahDavison
## 44 ScottHillis
## 45 SimonCowell
## 46 TanEeLyn
## 47 TheresePoletti
## 48 TimFarrand
## 49 ToddNissen
## 50 WilliamKazer
## total_correct_vector.predicted_per_author
## 1 0.49152542
## 2 0.06697819

```

## 3	NaN
## 4	0.21333333
## 5	0.72727273
## 6	NaN
## 7	NaN
## 8	0.05932203
## 9	NaN
## 10	1.00000000
## 11	0.87500000
## 12	1.00000000
## 13	0.50000000
## 14	NaN
## 15	0.21428571
## 16	0.28057554
## 17	1.00000000
## 18	1.00000000
## 19	NaN
## 20	0.75000000
## 21	0.66666667
## 22	0.50000000
## 23	NaN
## 24	1.00000000
## 25	NaN
## 26	0.16460905
## 27	0.07925801
## 28	0.79310345
## 29	1.00000000
## 30	0.60000000
## 31	NaN
## 32	NaN
## 33	1.00000000
## 34	0.50000000
## 35	NaN
## 36	0.62500000
## 37	NaN
## 38	0.43243243
## 39	1.00000000
## 40	0.95238095
## 41	0.43750000
## 42	1.00000000
## 43	NaN
## 44	NaN
## 45	NaN
## 46	NaN
## 47	0.50000000
## 48	0.26000000
## 49	0.00000000
## 50	NaN

We achieve a classification rate of 18.52%, which isn't superb. Moreover, there are 15 authors who we never predict to have written an article while there are 7 authors we predict to have written over 100 articles. In short, the results from Naive Bayes are inconsistent at best.

I will try a random forest model to see if I get better results.

```
TrainDataFrame = as.data.frame(Train_Matrix)
TestDataFrame = as.data.frame(Test_Matrix)

set.seed(722)
AuthorRandomForest = randomForest(x=TrainDataFrame,
y=as.factor(train_labels), ntree=50, mtry=30)

PredictedAuthor = predict(AuthorRandomForest, newdata = TestDataFrame)
```

I set the number of trees arbitrarily to be 50. The recommended number of variables to consider for categorical random forest problems is the square root of the number of predictor variables. In this case, that is the 1189 words remaining after the tokenization process. I round down from approximately 34 to arrive at 30 words. I then fit my random forest model to the test data set.

```
RandomForestResults = as.data.frame(table(PredictedAuthor, test_labels))
RandomForestResultsTable = table(PredictedAuthor, test_labels)
RF_total_correct_vector = rep(0, 50)
RF_predicted_per_author = rep(0, 50)
for (i in 1:length(RandomForestResultsTable[1,])){
  RF_total_correct_vector[i] = RandomForestResultsTable[i, i]
  RF_predicted_per_author[i] = sum(RandomForestResultsTable[i,])
}
```

```
RFCorrectByAuthor =  
data.frame(row.names(RandomForestResultsTable), RF_total_correct_vector)  
RFCorrectByAuthor
```

##	row.names.RandomForestResultsTable.	RF_total_correct_vector
## 1	AaronPressman	42
## 2	AlanCrosby	30
## 3	AlexanderSmith	19
## 4	BenjaminKangLim	16
## 5	BernardHickey	31
## 6	BradDorfman	29
## 7	DarrenSchuettler	14
## 8	DavidLawder	9
## 9	EdnaFernandes	19
## 10	EricAuchard	19
## 11	FumikoFujisaki	50
## 12	GrahamEarnshaw	43
## 13	HeatherScoffield	19
## 14	JaneMacartney	8

## 15	JanLopatka	32
## 16	JimGilchrist	50
## 17	JoeOrtiz	19
## 18	JohnMastrini	21
## 19	JonathanBirt	31
## 20	JoWinterbottom	37
## 21	KarlPenhaul	45
## 22	KeithWeir	34
## 23	KevinDrawbaugh	23
## 24	KevinMorrison	23
## 25	KirstinRidley	26
## 26	KouroshKarimkhany	34
## 27	LydiaZajc	32
## 28	LynneO'Donnell	40
## 29	LynnleyBrowning	49
## 30	MarcelMichelson	44
## 31	MarkBendeich	41
## 32	MartinWolk	22
## 33	MatthewBunce	46
## 34	MichaelConnor	27
## 35	MureDickie	16
## 36	NickLouth	40
## 37	PatriciaCommins	28
## 38	PeterHumphrey	29
## 39	PierreTran	21
## 40	RobinSidel	39
## 41	RogerFillion	39
## 42	SamuelPerry	22
## 43	SarahDavison	24
## 44	ScottHillis	18
## 45	SimonCowell	34
## 46	TanEeLyn	29
## 47	TheresePoletti	14
## 48	TimFarrand	23
## 49	ToddNissen	29
## 50	WilliamKazer	12

```

RFPredictedForAuthor = data.frame(row.names(RandomForestResultsTable),
RF_predicted_per_author)
RFSortedPredictedForAuthor = RFPredictedForAuthor[order(-
RFPredictedForAuthor$RF_predicted_per_author),]
RFSortedPredictedForAuthor

```

##	row.names.RandomForestResultsTable.	RF_predicted_per_author
## 46	TanEeLyn	82
## 30	MarcelMichelson	79
## 19	JonathanBirt	74
## 21	KarlPenhaul	70
## 49	ToddNissen	69
## 26	KouroshKarimkhany	68

## 4	BenjaminKangLim	65
## 15	JanLopatka	65
## 12	GrahamEarnshaw	63
## 16	JimGilchrist	63
## 38	PeterHumphrey	61
## 31	MarkBendeich	60
## 1	AaronPressman	59
## 13	HeatherScoffield	59
## 17	JoeOrtiz	57
## 34	MichaelConnor	57
## 36	NickLouth	57
## 11	FumikoFujisaki	56
## 7	DarrenSchuettler	55
## 37	PatriciaCommins	55
## 6	BradDorfman	54
## 29	LynnleyBrowning	54
## 33	MatthewBunce	53
## 40	RobinSidel	53
## 45	SimonCowell	52
## 5	BernardHickey	49
## 41	RogerFillion	49
## 44	ScottHillis	49
## 9	EdnaFernandes	48
## 20	JoWinterbottom	48
## 23	KevinDrawbaugh	48
## 22	KeithWeir	47
## 42	SamuelPerry	44
## 18	JohnMastrini	43
## 28	LynneO'Donnell	42
## 2	AlanCrosby	41
## 10	EricAuchard	40
## 43	SarahDavison	40
## 48	TimFarrand	39
## 24	KevinMorrison	37
## 35	MureDickie	37
## 27	LydiaZajc	33
## 3	AlexanderSmith	32
## 25	KirstinRidley	32
## 50	WilliamKazer	32
## 47	TheresePoletti	29
## 32	MartinWolk	28
## 14	JaneMacartney	27
## 39	PierreTran	27
## 8	DavidLawder	19

```
RFOverallClassificationRate = sum(RF_total_correct_vector)/2500
```

```
RFOverallClassificationRate
```

```
## [1] 0.5764
```

```
RFPrecisionRateByAuthor = data.frame(row.names(RandomForestResultsTable),  
RF_total_correct_vector/RF_predicted_per_author)  
RFPrecisionRateByAuthor
```

```
##      row.names.RandomForestResultsTable.  
## 1      AaronPressman  
## 2      AlanCrosby  
## 3      AlexanderSmith  
## 4      BenjaminKangLim  
## 5      BernardHickey  
## 6      BradDorfman  
## 7      DarrenSchuettler  
## 8      DavidLawder  
## 9      EdnaFernandes  
## 10     EricAuchard  
## 11     FumikoFujisaki  
## 12     GrahamEarnshaw  
## 13     HeatherScoffield  
## 14     JaneMacartney  
## 15     JanLopatka  
## 16     JimGilchrist  
## 17     JoeOrtiz  
## 18     JohnMastrini  
## 19     JonathanBirt  
## 20     JoWinterbottom  
## 21     KarlPenhaul  
## 22     KeithWeir  
## 23     KevinDrawbaugh  
## 24     KevinMorrison  
## 25     KirstinRidley  
## 26     KouroshKarimkhany  
## 27     LydiaZajc  
## 28     LynneO'Donnell  
## 29     LynnleyBrowning  
## 30     MarcelMichelson  
## 31     MarkBendeich  
## 32     MartinWolk  
## 33     MatthewBunce  
## 34     MichaelConnor  
## 35     MureDickie  
## 36     NickLouth  
## 37     PatriciaCommins  
## 38     PeterHumphrey  
## 39     PierreTran  
## 40     RobinSidel  
## 41     RogerFillion  
## 42     SamuelPerry  
## 43     SarahDavison  
## 44     ScottHillis  
## 45     SimonCowell
```

```
## 46          TanEeLyn
## 47      TheresePoletti
## 48          TimFarrand
## 49          ToddNissen
## 50      WilliamKazer
## RF_total_correct_vector.RF_predicted_per_author
## 1          0.7118644
## 2          0.7317073
## 3          0.5937500
## 4          0.2461538
## 5          0.6326531
## 6          0.5370370
## 7          0.2545455
## 8          0.4736842
## 9          0.3958333
## 10         0.4750000
## 11         0.8928571
## 12         0.6825397
## 13         0.3220339
## 14         0.2962963
## 15         0.4923077
## 16         0.7936508
## 17         0.3333333
## 18         0.4883721
## 19         0.4189189
## 20         0.7708333
## 21         0.6428571
## 22         0.7234043
## 23         0.4791667
## 24         0.6216216
## 25         0.8125000
## 26         0.5000000
## 27         0.9696970
## 28         0.9523810
## 29         0.9074074
## 30         0.5569620
## 31         0.6833333
## 32         0.7857143
## 33         0.8679245
## 34         0.4736842
## 35         0.4324324
## 36         0.7017544
## 37         0.5090909
## 38         0.4754098
## 39         0.7777778
## 40         0.7358491
## 41         0.7959184
## 42         0.5000000
## 43         0.6000000
## 44         0.3673469
```

```
## 45          0.6538462
## 46          0.3536585
## 47          0.4827586
## 48          0.5897436
## 49          0.4202899
## 50          0.3750000
```

I'm immediately much happier with my results. The overall classification rate jumps to 58%. Furthermore, there is a much better distribution of predicted number of articles for each author. This statistic ranges from 19 to 82 instead of 0 to 554 as in the Naive Bayes case.

Another nice thing about the Random Forest model is I can see which of the tokens were most important in fitting the model.

```
RFImportance = as.data.frame(importance(AuthorRandomForest))
RFImportanceDataFrame = data.frame(row.names(RFImportance), RFImportance)
SortRFImportance = RFImportanceDataFrame[order(-
RFImportanceDataFrame$MeanDecreaseGini),]
SortRFImportance[1:10,]
```

```
##          row.names.RFImportance. MeanDecreaseGini
## czech          czech          18.81620
## toronto        toronto          13.90261
## kong           kong           13.04951
## hong           hong           12.91905
## cargo          cargo          12.87677
## french         french          11.79951
## chinas         chinas          11.51090
## australian    australian       11.28318
## china          china          10.92082
## chinese        chinese          10.84070
```

Interestingly, many international words appear on the list. This leads me to believe that the authors we were most successfully able to classify write mostly about international items for the paper.

Association Rules

Practice with Association Rule Mining

First, I read in the data directly as transaction data and view a summary.

```
Groceries = read.transactions("../data/groceries.txt", format = "basket",
sep=",")
summary(Groceries)

## transactions as itemMatrix in sparse format with
## 9835 rows (elements/itemsets/transactions) and
## 169 columns (items) and a density of 0.02609146
```

```
##
## most frequent items:
##      whole milk other vegetables      rolls/buns      soda
##      2513      1903      1809      1715
##      yogurt      (Other)
##      1372      34055
##
## element (itemset/transaction) length distribution:
## sizes
##      1      2      3      4      5      6      7      8      9      10      11      12      13      14      15
## 2159 1643 1299 1005  855  645  545  438  350  246  182  117  78  77  55
##      16      17      18      19      20      21      22      23      24      26      27      28      29      32
##      46      29      14      14      9      11      4      6      1      1      1      1      3      1
##
##      Min. 1st Qu.  Median      Mean 3rd Qu.      Max.
##      1.000  2.000  3.000  4.409  6.000 32.000
##
## includes extended item information - examples:
##      labels
## 1 abrasive cleaner
## 2 artif. sweetener
## 3  baby cosmetics
```

Whole milk is the most popular item, followed by a generic vegetable category and a generic bread category.

Next, I will create association rules for these transactions using arbitrary cutoffs for support, confidence, and number of items allowed in a rule. I will explore the cutoffs in more detail next.

```
GroceriesRules <- apriori(Groceries, parameter=list(support=.005,
confidence=.5, maxlen=4))
##
## Parameter specification:
## confidence minval  smax  arem  aval originalSupport support minlen maxlen
##           0.5    0.1    1 none  FALSE             TRUE   0.005     1     4
## target  ext
## rules FALSE
##
## Algorithmic control:
## filter tree heap memopt load sort verbose
##    0.1 TRUE TRUE  FALSE TRUE    2    TRUE
##
## apriori - find association rules with the apriori algorithm
## version 4.21 (2004.05.09)      (c) 1996-2004  Christian Borgelt
## set item appearances ...[0 item(s)] done [0.00s].
## set transactions ...[169 item(s), 9835 transaction(s)] done [0.00s].
## sorting and recoding items ... [120 item(s)] done [0.00s].
## creating transaction tree ... done [0.00s].
## checking subsets of size 1 2 3 4 done [0.00s].
```

```
## writing ... [120 rule(s)] done [0.00s].  
## creating S4 object ... done [0.00s].
```

```
arules::inspect(GroceriesRules)
```

##	lhs	rhs	support
	confidence lift		
## 1	{baking powder}	=> {whole milk}	0.009252669
	0.5229885 2.046793		
## 2	{oil,		
##	other vegetables}	=> {whole milk}	0.005083884
	0.5102041 1.996760		
## 3	{onions,		
##	root vegetables}	=> {other vegetables}	0.005693950
	0.6021505 3.112008		
## 4	{onions,		
##	whole milk}	=> {other vegetables}	0.006609049
	0.5462185 2.822942		
## 5	{hygiene articles,		
##	other vegetables}	=> {whole milk}	0.005185562
	0.5425532 2.123363		
## 6	{other vegetables,		
##	sugar}	=> {whole milk}	0.006304016
	0.5849057 2.289115		
## 7	{long life bakery product,		
##	other vegetables}	=> {whole milk}	0.005693950
	0.5333333 2.087279		
## 8	{cream cheese ,		
##	yogurt}	=> {whole milk}	0.006609049
	0.5327869 2.085141		
## 9	{chicken,		
##	root vegetables}	=> {other vegetables}	0.005693950
	0.5233645 2.704829		
## 10	{chicken,		
##	root vegetables}	=> {whole milk}	0.005998983
	0.5514019 2.157993		
## 11	{chicken,		
##	rolls/buns}	=> {whole milk}	0.005287239
	0.5473684 2.142208		
## 12	{coffee,		
##	yogurt}	=> {whole milk}	0.005083884
	0.5208333 2.038359		
## 13	{frozen vegetables,		
##	root vegetables}	=> {other vegetables}	0.006100661
	0.5263158 2.720082		
## 14	{frozen vegetables,		
##	root vegetables}	=> {whole milk}	0.006202339
	0.5350877 2.094146		
## 15	{frozen vegetables,		
##	rolls/buns}	=> {whole milk}	0.005083884

0.5000000 1.956825		
## 16 {frozen vegetables,		
## other vegetables}	=> {whole milk}	0.009659380
0.5428571 2.124552		
## 17 {beef,		
## yogurt}	=> {whole milk}	0.006100661
0.5217391 2.041904		
## 18 {beef,		
## rolls/buns}	=> {whole milk}	0.006812405
0.5000000 1.956825		
## 19 {curd,		
## whipped/sour cream}	=> {whole milk}	0.005897306
0.5631068 2.203802		
## 20 {curd,		
## tropical fruit}	=> {yogurt}	0.005287239
0.5148515 3.690645		
## 21 {curd,		
## tropical fruit}	=> {other vegetables}	0.005287239
0.5148515 2.660833		
## 22 {curd,		
## tropical fruit}	=> {whole milk}	0.006507372
0.6336634 2.479936		
## 23 {curd,		
## root vegetables}	=> {other vegetables}	0.005490595
0.5046729 2.608228		
## 24 {curd,		
## root vegetables}	=> {whole milk}	0.006202339
0.5700935 2.231146		
## 25 {curd,		
## yogurt}	=> {whole milk}	0.010066090
0.5823529 2.279125		
## 26 {curd,		
## rolls/buns}	=> {whole milk}	0.005897306
0.5858586 2.292845		
## 27 {curd,		
## other vegetables}	=> {whole milk}	0.009862735
0.5739645 2.246296		
## 28 {pork,		
## root vegetables}	=> {other vegetables}	0.007015760
0.5149254 2.661214		
## 29 {pork,		
## root vegetables}	=> {whole milk}	0.006812405
0.5000000 1.956825		
## 30 {pork,		
## rolls/buns}	=> {whole milk}	0.006202339
0.5495495 2.150744		
## 31 {frankfurter,		
## tropical fruit}	=> {whole milk}	0.005185562
0.5483871 2.146195		
## 32 {frankfurter,		

## root vegetables} 0.5000000 1.956825	=> {whole milk}	0.005083884
## 33 {frankfurter, ## yogurt}	=> {whole milk}	0.006202339
0.5545455 2.170296		
## 34 {bottled beer, ## yogurt}	=> {whole milk}	0.005185562
0.5604396 2.193364		
## 35 {brown bread, ## tropical fruit}	=> {whole milk}	0.005693950
0.5333333 2.087279		
## 36 {brown bread, ## root vegetables}	=> {whole milk}	0.005693950
0.5600000 2.191643		
## 37 {brown bread, ## other vegetables}	=> {whole milk}	0.009354347
0.5000000 1.956825		
## 38 {domestic eggs, ## margarine}	=> {whole milk}	0.005185562
0.6219512 2.434099		
## 39 {margarine, ## root vegetables}	=> {other vegetables}	0.005897306
0.5321101 2.750028		
## 40 {margarine, ## rolls/buns}	=> {whole milk}	0.007930859
0.5379310 2.105273		
## 41 {butter, ## domestic eggs}	=> {whole milk}	0.005998983
0.6210526 2.430582		
## 42 {butter, ## whipped/sour cream}	=> {other vegetables}	0.005795628
0.5700000 2.945849		
## 43 {butter, ## whipped/sour cream}	=> {whole milk}	0.006710727
0.6600000 2.583008		
## 44 {butter, ## citrus fruit}	=> {whole milk}	0.005083884
0.5555556 2.174249		
## 45 {bottled water, ## butter}	=> {whole milk}	0.005388917
0.6022727 2.357084		
## 46 {butter, ## tropical fruit}	=> {other vegetables}	0.005490595
0.5510204 2.847759		
## 47 {butter, ## tropical fruit}	=> {whole milk}	0.006202339
0.6224490 2.436047		
## 48 {butter, ## root vegetables}	=> {other vegetables}	0.006609049
0.5118110 2.645119		

## 49 {butter, ## root vegetables} 0.6377953 2.496107	=> {whole milk}	0.008235892
## 50 {butter, ## yogurt} 0.6388889 2.500387	=> {whole milk}	0.009354347
## 51 {butter, ## other vegetables} 0.5736041 2.244885	=> {whole milk}	0.011489578
## 52 {newspapers, ## root vegetables} 0.5221239 2.698417	=> {other vegetables}	0.005998983
## 53 {newspapers, ## root vegetables} 0.5044248 1.974142	=> {whole milk}	0.005795628
## 54 {domestic eggs, ## whipped/sour cream} 0.5102041 2.636814	=> {other vegetables}	0.005083884
## 55 {domestic eggs, ## whipped/sour cream} 0.5714286 2.236371	=> {whole milk}	0.005693950
## 56 {domestic eggs, ## pip fruit} 0.6235294 2.440275	=> {whole milk}	0.005388917
## 57 {citrus fruit, ## domestic eggs} 0.5490196 2.148670	=> {whole milk}	0.005693950
## 58 {domestic eggs, ## tropical fruit} 0.6071429 2.376144	=> {whole milk}	0.006914082
## 59 {domestic eggs, ## root vegetables} 0.5106383 2.639058	=> {other vegetables}	0.007320793
## 60 {domestic eggs, ## root vegetables} 0.5957447 2.331536	=> {whole milk}	0.008540925
## 61 {domestic eggs, ## yogurt} 0.5390071 2.109485	=> {whole milk}	0.007727504
## 62 {domestic eggs, ## other vegetables} 0.5525114 2.162336	=> {whole milk}	0.012302999
## 63 {fruit/vegetable juice, ## root vegetables} 0.5508475 2.846865	=> {other vegetables}	0.006609049
## 64 {fruit/vegetable juice, ## root vegetables} 0.5423729 2.122657	=> {whole milk}	0.006507372
## 65 {fruit/vegetable juice, ## yogurt}	=> {whole milk}	0.009456024

0.5054348 1.978094		
## 66 {pip fruit,		
## whipped/sour cream}	=> {other vegetables}	0.005592272
0.6043956 3.123610		
## 67 {pip fruit,		
## whipped/sour cream}	=> {whole milk}	0.005998983
0.6483516 2.537421		
## 68 {citrus fruit,		
## whipped/sour cream}	=> {other vegetables}	0.005693950
0.5233645 2.704829		
## 69 {citrus fruit,		
## whipped/sour cream}	=> {whole milk}	0.006304016
0.5794393 2.267722		
## 70 {sausage,		
## whipped/sour cream}	=> {whole milk}	0.005083884
0.5617978 2.198679		
## 71 {tropical fruit,		
## whipped/sour cream}	=> {other vegetables}	0.007829181
0.5661765 2.926088		
## 72 {tropical fruit,		
## whipped/sour cream}	=> {whole milk}	0.007930859
0.5735294 2.244593		
## 73 {root vegetables,		
## whipped/sour cream}	=> {other vegetables}	0.008540925
0.5000000 2.584078		
## 74 {root vegetables,		
## whipped/sour cream}	=> {whole milk}	0.009456024
0.5535714 2.166484		
## 75 {whipped/sour cream,		
## yogurt}	=> {whole milk}	0.010879512
0.5245098 2.052747		
## 76 {rolls/buns,		
## whipped/sour cream}	=> {whole milk}	0.007829181
0.5347222 2.092715		
## 77 {other vegetables,		
## whipped/sour cream}	=> {whole milk}	0.014641586
0.5070423 1.984385		
## 78 {pip fruit,		
## sausage}	=> {whole milk}	0.005592272
0.5188679 2.030667		
## 79 {pip fruit,		
## root vegetables}	=> {other vegetables}	0.008134215
0.5228758 2.702304		
## 80 {pip fruit,		
## root vegetables}	=> {whole milk}	0.008947636
0.5751634 2.250988		
## 81 {pip fruit,		
## yogurt}	=> {whole milk}	0.009557702
0.5310734 2.078435		
## 82 {other vegetables,		

## pip fruit} 0.5175097 2.025351	=> {whole milk}	0.013523132
## 83 {pastry, ## tropical fruit} 0.5076923 1.986930	=> {whole milk}	0.006710727
## 84 {pastry, ## root vegetables} 0.5370370 2.775491	=> {other vegetables}	0.005897306
## 85 {pastry, ## root vegetables} 0.5185185 2.029299	=> {whole milk}	0.005693950
## 86 {pastry, ## yogurt} 0.5172414 2.024301	=> {whole milk}	0.009150991
## 87 {citrus fruit, ## root vegetables} 0.5862069 3.029608	=> {other vegetables}	0.010371124
## 88 {citrus fruit, ## root vegetables} 0.5172414 2.024301	=> {whole milk}	0.009150991
## 89 {root vegetables, ## shopping bags} 0.5158730 2.666112	=> {other vegetables}	0.006609049
## 90 {sausage, ## tropical fruit} 0.5182482 2.028241	=> {whole milk}	0.007219115
## 91 {root vegetables, ## sausage} 0.5170068 2.023383	=> {whole milk}	0.007727504
## 92 {root vegetables, ## tropical fruit} 0.5845411 3.020999	=> {other vegetables}	0.012302999
## 93 {root vegetables, ## tropical fruit} 0.5700483 2.230969	=> {whole milk}	0.011997966
## 94 {tropical fruit, ## yogurt} 0.5173611 2.024770	=> {whole milk}	0.015149975
## 95 {root vegetables, ## yogurt} 0.5000000 2.584078	=> {other vegetables}	0.012913066
## 96 {root vegetables, ## yogurt} 0.5629921 2.203354	=> {whole milk}	0.014539908
## 97 {rolls/buns, ## root vegetables} 0.5020921 2.594890	=> {other vegetables}	0.012201322
## 98 {rolls/buns, ## root vegetables} 0.5230126 2.046888	=> {whole milk}	0.012709710

## 99 {other vegetables, ## yogurt} 0.5128806 2.007235	=> {whole milk}	0.022267412
## 100 {fruit/vegetable juice, ## other vegetables, ## yogurt} 0.6172840 2.415833	=> {whole milk}	0.005083884
## 101 {fruit/vegetable juice, ## whole milk, ## yogurt} 0.5376344 2.778578	=> {other vegetables}	0.005083884
## 102 {other vegetables, ## root vegetables, ## whipped/sour cream} 0.6071429 2.376144	=> {whole milk}	0.005185562
## 103 {root vegetables, ## whipped/sour cream, ## whole milk} 0.5483871 2.834150	=> {other vegetables}	0.005185562
## 104 {other vegetables, ## whipped/sour cream, ## yogurt} 0.5500000 2.152507	=> {whole milk}	0.005592272
## 105 {whipped/sour cream, ## whole milk, ## yogurt} 0.5140187 2.656529	=> {other vegetables}	0.005592272
## 106 {other vegetables, ## pip fruit, ## root vegetables} 0.6750000 2.641713	=> {whole milk}	0.005490595
## 107 {pip fruit, ## root vegetables, ## whole milk} 0.6136364 3.171368	=> {other vegetables}	0.005490595
## 108 {other vegetables, ## pip fruit, ## yogurt} 0.6250000 2.446031	=> {whole milk}	0.005083884
## 109 {pip fruit, ## whole milk, ## yogurt} 0.5319149 2.749019	=> {other vegetables}	0.005083884
## 110 {citrus fruit, ## other vegetables, ## root vegetables} 0.5588235 2.187039	=> {whole milk}	0.005795628
## 111 {citrus fruit, ## root vegetables, ## whole milk}	=> {other vegetables}	0.005795628

```

0.6333333 3.273165
## 112 {root vegetables,
##     tropical fruit,
##     yogurt}          => {whole milk}          0.005693950
0.7000000 2.739554
## 113 {other vegetables,
##     root vegetables,
##     tropical fruit}   => {whole milk}          0.007015760
0.5702479 2.231750
## 114 {root vegetables,
##     tropical fruit,
##     whole milk}       => {other vegetables} 0.007015760
0.5847458 3.022057
## 115 {other vegetables,
##     tropical fruit,
##     yogurt}          => {whole milk}          0.007625826
0.6198347 2.425816
## 116 {tropical fruit,
##     whole milk,
##     yogurt}          => {other vegetables} 0.007625826
0.5033557 2.601421
## 117 {other vegetables,
##     root vegetables,
##     yogurt}          => {whole milk}          0.007829181
0.6062992 2.372842
## 118 {root vegetables,
##     whole milk,
##     yogurt}          => {other vegetables} 0.007829181
0.5384615 2.782853
## 119 {other vegetables,
##     rolls/buns,
##     root vegetables}  => {whole milk}          0.006202339
0.5083333 1.989438
## 120 {other vegetables,
##     rolls/buns,
##     yogurt}          => {whole milk}          0.005998983
0.5221239 2.043410

```

These cutoffs yield 120 rules. While this is a lot to parse, I notice that a lot of the rules are used to predict when someone will buy whole milk. Since whole milk is the most common item, I want to look at lift in an attempt to "normalize" against how frequent an item is.

```

arules::inspect(subset(GroceriesRules, subset=lift > 3))

##   lhs                                rhs          support confidence
##   lift
## 1 {onions,
##    root vegetables}    => {other vegetables} 0.005693950  0.6021505
3.112008
## 2 {curd,

```

```

##   tropical fruit}      => {yogurt}          0.005287239  0.5148515
3.690645
## 3 {pip fruit,
##   whipped/sour cream} => {other vegetables} 0.005592272  0.6043956
3.123610
## 4 {citrus fruit,
##   root vegetables}    => {other vegetables} 0.010371124  0.5862069
3.029608
## 5 {root vegetables,
##   tropical fruit}     => {other vegetables} 0.012302999  0.5845411
3.020999
## 6 {pip fruit,
##   root vegetables,
##   whole milk}         => {other vegetables} 0.005490595  0.6136364
3.171368
## 7 {citrus fruit,
##   root vegetables,
##   whole milk}         => {other vegetables} 0.005795628  0.6333333
3.273165
## 8 {root vegetables,
##   tropical fruit,
##   whole milk}         => {other vegetables} 0.007015760  0.5847458
3.022057

```

There are 8 rules with a lift of more than 3, indicating they have predictive power beyond just predicting a popular item will be in a basket. These rules don't predict purchasing whole milk, but they do predict purchasing other vegetables, which was the second most common item. Many of the rules make sense intuitively. For instance, people who buy citrus fruit and root vegetables are already shopping for produce, so buying other vegetables isn't a stretch of the imagination.

Let's now look at rules which occur in more than 1.5% of all transactions.

```

arules::inspect(subset(GroceriesRules, subset=support > 0.015))
##   lhs                rhs                support confidence    lift
## 1 {tropical fruit,
##   yogurt}            => {whole milk} 0.01514997  0.5173611 2.024770
## 2 {other vegetables,
##   yogurt}            => {whole milk} 0.02226741  0.5128806 2.007235

```

There are only two rules which occur more than 1.5% of the time, and both are used to predict whole milk. From this, we can see that if a customer buys yogurt and produce, they are also very likely to buy whole milk.

Finally, I want to look at the cases where the predicted item occurs most frequently with the predictor items.

```

arules::inspect(subset(GroceriesRules, subset=support > .003 & confidence > 0.65))

```


##	lhs	rhs	support	confidence	lift
## 1	{butter,				
##	whipped/sour cream}	=> {whole milk}	0.006710727	0.660	2.583008
## 2	{other vegetables,				
##	pip fruit,				
##	root vegetables}	=> {whole milk}	0.005490595	0.675	2.641713
## 3	{root vegetables,				
##	tropical fruit,				
##	yogurt}	=> {whole milk}	0.005693950	0.700	2.739554

All of the predicted items in this case involve whole milk. Again, produce and other dairy products show up. These rules all have high lift as well.